

Exercise Sheet 1

No Fear of Numbers: Introduction to Quantitative Data Analysis in R

Dr. Susanne de Vogel & Dr. Maryam Movahedifar, Data Science Center, University of Bremen

11 March 2026

Research Question: Is there a Life beyond the PhD?

In this workshop, we will work with one coherent example throughout all four parts. We focus on two outcomes of interest:

- `bpsy01`: overall life satisfaction (10 point scale)
- `blcd06`: Children (yes/no)

We will examine how these outcomes are related to different aspects of doctoral researchers' lives and backgrounds:

PhD and Work conditions

- `adbi01`: Status of the doctorate
- `adbi15`: Discipline
- `bcd17`: Emotional Support during PhD
- `bwd12`: Perceived scientific pressure
- `bemp81`: Monthly gross income

Attitudes and Well-Being

- `blcd12`: Satisfaction with work-life balance
- `apar15b`: Relationship with parents
- `bpsy05`: Self-Efficacy

Demographics

- `adem01`: Gender
- `adem02`: Age in years
- `apar10`: Highest vocational degree of parents
- `adem03`: Country of birth
- `blcd01`: Relationship status

Exploring datasets

In this exercise, we will use R to explore the data frame. The goal is to learn different ways to *examine the data frame and its variables* as a *preparation for analysis*. You will learn how to

- use different functions to inspect the data frame's structure
- examine variable types and classes
- check labels and missing values

0 Set-up: Installing and loading packages, loading the data frame

Okaaaaaay, let's go!

0.1. Install `tidyverse` and `haven` if not already installed.

When using the packages for the *first time*, you could simply install it with `install.packages()`, e.g. `install.packages("tidyverse")`.

Alternatively, you can use a *conditional check* like

```
if (!requireNamespace("tidyverse", quietly = TRUE)) install.packages("tidyverse")
```

`requireNamespace("tidyverse", quietly = TRUE)` returns TRUE if the package is available. The `!` means “not”, so the condition becomes TRUE only when the package is missing. In that case, `install.packages("tidyverse")` is executed. This checks automatically whether the package is already installed.

To do so, start your script with the following code:

```
# Install tidyverse
if (!requireNamespace("tidyverse", quietly = TRUE)) # quietly = TRUE means that R runs the function
  ↪ silently
  install.packages("tidyverse")

# Install haven
if (!requireNamespace("haven", quietly = TRUE))
  install.packages("haven")
```

0.2. Use the `library()` function to load *tidyverse* and *haven* so that we can work with them.

Hint: Use the same name of the package, but *without* quotation marks.

Solution:

```
# Load the packages (every time you start R)
library(tidyverse)
library(haven)
```

Note: What happens when you load `tidyverse`?

- You might see a **Warning** like: “*Package was built under R version 4.x.y*”. This is **not an error**. It only means the package was created with a newer R version. In almost all cases, everything will work fine.
- R shows which **core tidyverse packages** are attached automatically: `ggplot2`, `dplyr`, `tidyr`, `readr`, `purrr`, `tibble`, `stringr`, `forcats` etc.

- Under **Conflicts**, R warns you if functions from different packages have the same name (e.g. `filter()` exists in both `dplyr` and `stats`). By default, the tidyverse version is used. If you want the other one, call it explicitly, e.g. `stats::filter()`.

0.3. Define the *Exercises* folder as your working directory with the function `setwd("path")`.

Hint: Note that R uses a slash/ instead of a backslash .

Solution:

```
# define working directory
setwd(
  ↪ "C:/Users/Susanne/Nextcloud/share_DSC/003_Trainings/001_Workshops/2026/2026-03-11_SdV_MM_Quantitative_Data_Anal
  ↪ )
```

0.4. Load the SPSS file “01_qa_exploring_data.sav” as a data frame using the `read_sav("dataset file name")` function and call this data frame `mydata1`.

Solution:

```
# Load the Nacaps dataset and name this data frame "mydata1"
mydata1 <- read_sav("01_qa_exploring_data.sav")
```

Exercise 1: Exploring the structure of your data frame

To check whether the data are in a tidy structure, **tidyverse** provides a range of functions that allow us to explore the data frame and its variables.

1.1. On the right side, in the Environment pane, the data frame `mydata1` should now appear under *Data*. What kind of information about the data frame do you find here?

Solution:

Firstly, there is the **number of observations**, e.g. the number of participants in your data frame. This corresponds to the number of rows in your data frame. Secondly, the **number of variables**, e.g. the attributes in your data frame. This corresponds to the number of columns in your data frame.

1.2. Use the function `glimpse(mydata1)` to take a first look at your data frame. What kind of information do you get from the output?

Solution:

```
# Exploring the structure of the data frame
glimpse(mydata1)

## Rows: 2,354
## Columns: 16
## $ pid      <dbl> 2, 15, 39, 55, 75, 81, 90, 104, 116, 117, 137, 165, 166, 167, ~
## $ adbi15   <dbl+lbl> 7, 4, 1, 3, 5, 1, 4, 4, 3, 4, 1, 4, NA, 7, 5~
## $ adem01   <dbl+lbl> 2, 2, 1, 1, 1, 2, 2, 2, 2, 1, 2, 2, 2, 2, 1, 1, 3, 2, 2, 1~
## $ adem02   <dbl> 27, 25, 41, 30, 26, 25, 27, 28, 34, 29, 34, 27, 42, 29, 31, 27~
## $ adem03   <dbl+lbl> 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1~
## $ apar10   <dbl+lbl> 3, 1, 2, 3, 3, 1, 3, 2, 3, 2, 2, 2, 1, 2, 1, 3, 1, 1, 2, 3~
## $ apar15b  <dbl+lbl> 3, 4, 2, 3, 3, 3, 4, 3, 3, 3, 3, 4, 3, 3, 4, 3, 3, 4, 3~
## $ bdbi01   <dbl+lbl> 1, 1, 1, 1, 3, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1~
## $ bdc17    <dbl> 3.000000, 4.333333, 4.000000, 4.666667, 1.000000, 3.333333, 4.~
```

```
## $ bemp81 <dbl> 3076, 0, 0, 0, 985, 1762, 4266, 4381, NA, NA, 2592, 660, 4285, ~
## $ bdwr12 <dbl> 2.50, 2.50, 2.00, 2.50, 1.00, 3.50, 2.50, 2.50, NA, NA, 4.25, ~
## $ bpsy01 <dbl+lbl> 5, 7, 7, 8, 5, 5, 7, 7, NA, NA, 8, 5, 8, 9, 8~
## $ blcd01 <dbl+lbl> 2, 2, 1, 1, 1, 2, 2, 2, NA, NA, 1, 2, 1, 1, 1~
## $ blcd06 <dbl+lbl> 2, 2, 2, 2, 2, 2, 2, 2, NA, NA, 1, 2, 1, 2, 1~
## $ blcd12 <dbl+lbl> 8, 3, 8, 8, 5, 10, 3, 5, NA, NA, 3, 7, 9, 9, 10~
## $ bpsy05 <dbl> 5.000000, 4.000000, 4.333333, 3.666667, 4.666667, 5.000000, 4.~
```

The function `glimpse()` provides a compact overview of the data frame. It shows

- how many rows and columns the data frame has
- the name of each variable
- the order of the variables
- the data type (e.g., character, double, integer, factor or `dbl+lbl`)
- a preview of the first few values

1.3. Execute the functions `head(mydata1)` and `tail(mydata1)`. What is the difference to `glimpse()`? Which information do you get?

Solution:

```
# display the first six rows of the dataset
head(mydata1)
```

```
## # A tibble: 6 x 16
##   pid adbi15      adem01 adem02 adem03 apar10 apar15b bdbi01 bdcd17 bemp81
##   <dbl> <dbl+lbl> <dbl+lbl> <dbl> <dbl+lbl> <dbl+lbl> <dbl+lbl> <dbl+lbl> <dbl> <dbl>
## 1     2 7 [enginee~ 2 [mal~      27 1 [in ~ 3 [Oth~ 3 [wel~ 1 [I a~      3      3076
## 2    15 4 [mathema~ 2 [mal~      25 1 [in ~ 1 [PhD~ 4 [ver~ 1 [I a~      4.33      0
## 3    39 1 [humanit~ 1 [fem~      41 2 [in ~ 2 [Bac~ 2 [not~ 1 [I a~      4      0
## 4    55 3 [law, ec~ 1 [fem~      30 1 [in ~ 3 [Oth~ 3 [wel~ 1 [I a~      4.67      0
## 5    75 5 [human m~ 1 [fem~      26 1 [in ~ 3 [Oth~ 3 [wel~ 3 [I h~      1      985
## 6    81 1 [humanit~ 2 [mal~      25 1 [in ~ 1 [PhD~ 3 [wel~ 1 [I a~      3.33     1762
## # i 6 more variables: bdwr12 <dbl>, bpsy01 <dbl+lbl>, blcd01 <dbl+lbl>,
## #   blcd06 <dbl+lbl>, blcd12 <dbl+lbl>, bpsy05 <dbl>
```

```
# display the last six rows of the dataset
tail(mydata1)
```

```
## # A tibble: 6 x 16
##   pid adbi15      adem01 adem02 adem03 apar10 apar15b bdbi01 bdcd17 bemp81
##   <dbl> <dbl+lbl> <dbl+lbl> <dbl> <dbl+lbl> <dbl+lbl> <dbl+lbl> <dbl+lbl> <dbl> <dbl>
## 1 28277 3 [law, ec~ 1 [fem~      29 1 [in ~ 2 [Bac~ 2 [not~ 1 [I a~      4      3074
## 2 28301 5 [human m~ 2 [mal~      30 2 [in ~ 1 [PhD~ 4 [ver~ 1 [I a~      5      2015
## 3 28335 4 [mathema~ 2 [mal~      39 1 [in ~ 1 [PhD~ 3 [wel~ 2 [I h~      2.67     5897
## 4 28337 3 [law, ec~ 2 [mal~      31 1 [in ~ 2 [Bac~ 4 [ver~ 1 [I a~      4.33      50
## 5 28341 1 [humanit~ 1 [fem~      29 1 [in ~ 1 [PhD~ 4 [ver~ 1 [I a~      4.67     2182
## 6 28342 7 [enginee~ 1 [fem~      31 1 [in ~ 1 [PhD~ 4 [ver~ 1 [I a~      5      NA
## # i 6 more variables: bdwr12 <dbl>, bpsy01 <dbl+lbl>, blcd01 <dbl+lbl>,
## #   blcd06 <dbl+lbl>, blcd12 <dbl+lbl>, bpsy05 <dbl>
```

`head(mydata1)` shows the first few rows (by default 6). `tail(mydata1)` shows the last few rows (by default 6).

Tipp: If you want to display more or less rows than the first/last six rows, you can specify it using e.g. `head(mydata1, 10)` for the first ten rows.

With these functions, you see *actual records from your dataset*, row by row, across all variables.

1.4. Inspect the data frame in the RStudio Data Viewer using the function `View(mydata1)`. What can you see here?

Be careful: R is case sensitive, so use V with a capital letter!

Solution:

```
View(mydata1)
```

The function `View()` opens the dataset in a spreadsheet-like window inside RStudio. You can

- investigate the variables (columns), cases (rows) and values (cells) to check for a tidy data structure.
- check variable names, labels and variable orders.
- scroll through rows and columns, sort variables, and visually inspect values. This is helpful for spotting obvious data entry errors, strange characters, or missing values.

Note: `View()` is only for looking at the data, **not for editing**. Any changes you make inside the Data Viewer will not be saved to your dataset. If you want to clean or transform the data, you need to use functions like `mutate()`, `filter()`, or `select()` in your code.

Exercise 2: Inspecting categorical variables

To understand what kind of variables we are dealing with, we now inspect a few example variables in more detail.

Your task is to find out which type and class of variable we have, which labels are available, whether there are any missings (NAs).

To call not a whole data frame, but a *specific variable* in that data frame, write `mydata1$adem01`. You can read this as: “the variable `adem01` from the data frame `mydata1`”:

- `mydata1` = name of the data frame (the table)
- `$` = “take this specific column from the table”
- `adem01` = name of the variable / column

2.1. Use the functions `typeof()` and `class()` to inspect the variable `adem01`. What type and class is it in R?

Solution:

```
# check variable type
typeof(mydata1$adem01)
```

```
## [1] "double"
```

```
# check variable class
class(mydata1$adem01)
```

```
## [1] "haven_labelled" "vctrs_vctr"      "double"
```

The output shows that the variable `adem01` is stored internally as a numeric (double) vector, but it has the additional class “haven_labelled”, meaning that it contains value labels (e.g., imported from SPSS, Stata, or SAS) while still being treated as a numeric variable in R.

2.2. Use `attr()` with the options "label" and "labels" to check whether a variable label and value labels are stored for `adem01`.

Code it like this:

```
attr(mydata1$adem01, "label")      # variable label
```

```
## [1] "Gender"
```

```
attr(mydata1$adem01, "labels")     # value labels
```

```
## female   male   other
##        1      2      3
```

The variable label is *Gender*, the values 1 to 3 are labelled 1 = *female*, 2 = *male* , 3 = *other*

2.3. Create a simple frequency table for `adem01` using `count()`.

Use the pipe operator `%>%` to first pass the whole data frame into `count()`, e.g. `mydata1 %>% count(adem01)`, because `count()` needs a data frame as input and cannot be applied directly to a single variable like `mydata1$adem01`.

Code it like this:

```
# Take the dataset mydata1 and then count the occurrences of each value of adem01.
mydata1 %>%
  count(adem01)
```

```
## # A tibble: 4 x 2
##   adem01      n
##   <dbl> <dbl>
## 1 1 [female] 1153
## 2 2 [male]   1068
## 3 3 [other]   132
## 4 NA         1
```

2.4. Check whether there are any missing values for `adem01` using the `sum(is.na())` function.

Solution:

```
# 3) Check for missing values
sum(is.na(mydata1$adem01))
```

```
## [1] 1
```

There is one missing value.

Exercise 3: Inspecting numeric variables.

3.1. Use the functions `typeof()` and `class()` to check the type and class of the variable `adem02`.

Solution:

```
# check variable type
typeof(mydata1$adem02)
```

```
## [1] "double"
```

```
# check variable class
class(mydata1$adem02)
```

```
## [1] "numeric"
```

The output indicates that the variable `adem02` is stored as a numeric (double) vector and has the standard R class “numeric”. This means it is a regular numeric variable without additional attributes such as value labels.

3.2. Use `attr()` to check the variable and value labels. Anything noticeable?

Solution:

```
attr(mydata1$adem02, "label")      # variable label
```

```
## [1] "Age in years 2019"
```

```
attr(mydata1$adem02, "labels")     # value labels
```

```
## NULL
```

We see that there are no value labels for numeric variables, as their values are self-explanatory.

3.3. Use `summary()` and `count()` (remember to write `count()` as a pipe) to get a quick overview of the range of ages.

Solution:

```
# quick numeric summary
summary(mydata1$adem02)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's
##    22.00  28.00   30.00   31.45  33.00   68.00    15
```

```
# Take the dataset mydata1 and then count the occurrences of each value of adem02.
mydata1 %>%
  count(adem02)
```

```
## # A tibble: 47 x 2
##   adem02      n
##   <dbl> <int>
## 1     22     25
## 2     23     15
## 3     24     19
## 4     25    113
## 5     26    113
## 6     27    236
## 7     28    215
## 8     29    275
## 9     30    288
## 10    31    195
## # i 37 more rows
```

There are 15 missing values.

Take Home Checklist: Exploring a new dataset in R

Step	What to do when you open a new dataset in R	Useful functions / questions
1	Check the overall structure	<code>glimpse()</code> , <code>head()</code> , <code>tail()</code> , <code>View()</code>
2	Understand key variables	<code>typeof()</code> , <code>class()</code> , <code>attr(..., "label")</code> , <code>attr(..., "labels")</code>
3	Look at typical value ranges	<code>summary()</code> , <code>count()</code>
4	Identify missing and special codes	<code>is.na()</code> , unusual values like -996, -998, -989
5	Take notes	Which variables look promising? Which ones clearly need cleaning?